



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

**Batch: A1      Roll No.:**

**16010120015 YASH**

**16010120006 DIKSHITA**

**16010120018 JINAY**

**Experiment / assignment / tutorial No. 7**

**Grade: AA / AB / BB / BC / CC / CD / DD**

**Title: Implementing indexing and query processing**

**Objective:** To understand Query Processing and implement indexing to improve query execution plans

**Expected Outcome of Experiment:**

CO 3: Use SQL for relational database creation , maintenance and query processing

**Books/ Journals/ Websites referred:**

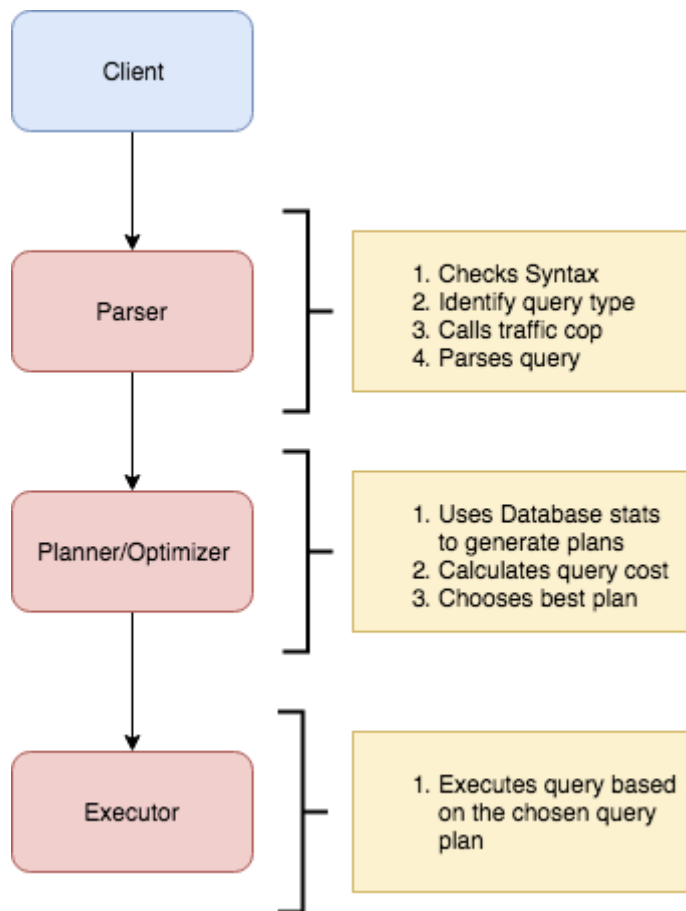
1. Dr. P.S. Deshpande, SQL and PL/SQL for Oracle 10g.Black book, Dreamtech Press
2. www.db-book.com
3. Korth, Silberchatz, Sudarshan : “Database Systems Concept”, 5<sup>th</sup> Edition , McGraw Hill
4. Elmasri and Navathe,”Fundamentals of database Systems”, 4<sup>th</sup> Edition,PEARSON Education.

**Resources used:** PostgreSQL

**Theory**



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)



Postgres query life cycle

### Planner and Executor

The planner receives a query tree from the rewriter and generates a (query) plan tree that can be processed by the executor most effectively.

The planner in PostgreSQL is based on pure cost-based optimization;

EXPLAIN command:

#### EXPLAIN command:

This command displays the execution plan that the PostgreSQL planner generates for the supplied statement. The execution plan shows how the table(s) referenced by the statement will be scanned — by plain sequential scan, index scan, etc. — and if multiple tables are referenced, what join algorithms will be used to bring together the required rows from each input table.

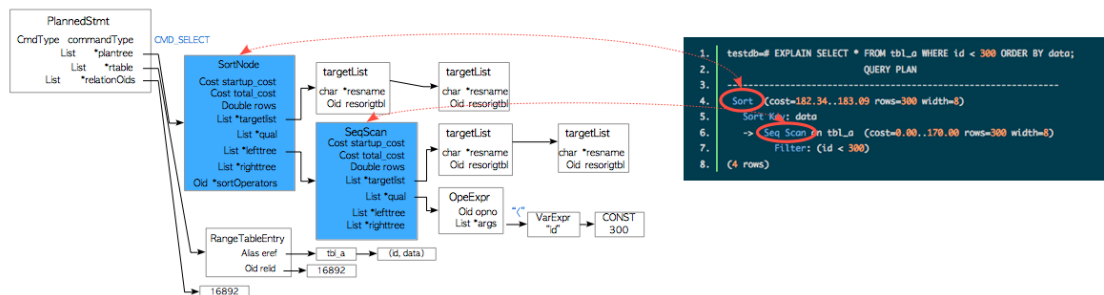


**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

As in the other RDBMS, the EXPLAIN command in PostgreSQL displays the plan tree itself. A specific example is shown below.

1. testdb=# EXPLAIN SELECT \* FROM tbl\_a WHERE id < 300 ORDER BY data;
2. QUERY PLAN
3. -----
4. Sort (cost=182.34..183.09 rows=300 width=8)
5. Sort Key: data
6. -> Seq Scan on tbl\_a (cost=0.00..170.00 rows=300 width=8)
7. Filter: (id < 300)
8. (4 rows)

**A simple plan tree and the relationship between the plan tree and the result of the EXPLAIN command.**



## Nodes

The first thing to understand is that each indented block with a preceding “->” (along with the top line) is called a node. A node is a logical unit of work (a “step” if you will) with an associated cost and execution time. The costs and times presented at each node are cumulative and roll up all child nodes.

**Cost:** It is not the time but a concept designed to estimate the cost of an operation. The first number is start up cost (cost to retrieve first record) and the second number is the cost incurred to process entire node (total cost from start to finish).



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

Cost is a combination of 5 work components used to estimate the work required: sequential fetch, non-sequential (random) fetch, processing of row, processing operator (function), and processing index entry.

**Rows** are the approximate number of rows returned when a specified operation is performed.

(In the case of select with where clause rows returned is

Rows = cardinality of relation \* selectivity )

**Width** is an average size of one row in bytes.

**Explain Analyze command:**

The EXPLAIN ANALYZE option causes the statement to be actually executed, not only planned. Then actual run time statistics are added to the display, including the total elapsed time expended within each plan node (in milliseconds) and the total number of rows it actually returned. This is useful for seeing whether the planner's estimates are close to reality.

```
EXPLAIN (ANALYZE) SELECT * FROM foo;
```

**QUERY PLAN**

— Seq Scan on foo (cost=0.00..18334.10 rows=1000010 width=37) (actual time=0.012..61.524 rows=1000010 loops=1)  
Total runtime: 90.944 ms  
(2 rows)

The command displays the following additional parameters:

- **actual time** is the actual time in milliseconds spent to get the first row and all rows, respectively.
- **rows** is the actual number of rows received with Seq Scan.
- **loops** is the number of times the Seq Scan operation had to be performed.
- **Total runtime** is the total time of query execution.

Query plans for select with where clause can be sequential scan, Index Scan, Index only Scan, Bitmap Index Scan etc.

Query plans for joins are Nested loop join, Hash join, Merge join etc.



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

### Implementation Screenshots:

Demonstrate query processing for the following types of query on your database

#### a. Simple select query

The screenshot shows the PostgreSQL Query Editor interface. The query entered is: `explain analyse select * from products;`. The 'Data Output' tab is selected, displaying the 'QUERY PLAN' for the query. The plan shows a 'Seq Scan on products' with a cost of 0.00..12.40, 240 rows, and a width of 312. The actual execution time was 0.286..0.289 seconds, with 5 rows returned in 1 loop. The planning time was 0.211 ms and the execution time was 2.455 ms.

Step	Operation	Cost	Rows	Width	Actual Time	Actual Rows	Loops
1	Seq Scan on products	0.00..12.40	240	312	0.286..0.289	5	1
2	Planning Time				0.211 ms		
3	Execution Time				2.455 ms		

#### b. Select query with where clause

The screenshot shows the PostgreSQL Query Editor interface. The query entered is: `explain analyse select * from products where mrp>700;`. The 'Data Output' tab is selected, displaying the 'QUERY PLAN' for the query. The plan shows a 'Seq Scan on products' with a cost of 0.00..13.00, 80 rows, and a width of 312. The actual execution time was 0.030..0.032 seconds, with 3 rows returned in 1 loop. The planning time was 0.162 ms and the execution time was 0.057 ms. The plan also indicates that 2 rows were removed by the filter.

Step	Operation	Cost	Rows	Width	Actual Time	Actual Rows	Loops
1	Seq Scan on products	0.00..13.00	80	312	0.030..0.032	3	1
2	Filter: (mrp > 700)						
3	Rows Removed by Filter		2				
4	Planning Time				0.162 ms		
5	Execution Time				0.057 ms		

#### c. Select query with order by query



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

postgres/postgres@PostgreSQL 13

Query Editor Query History

```
1 explain analyse select * from products where mrp>700 order by product_id;
```

Data Output Explain Messages Notifications

QUERY PLAN

text

1	Sort (cost=15.53..15.73 rows=80 width=312) (actual time=0.052..0.053 rows=3 loops=1)
2	Sort Key: product_id
3	Sort Method: quicksort Memory: 25kB
4	-> Seq Scan on products (cost=0.00..13.00 rows=80 width=312) (actual time=0.023..0.024 rows=3 loops=1)
5	Filter: (mrp > 700)
6	Rows Removed by Filter: 2
7	Planning Time: 0.191 ms
8	Execution Time: 0.081 ms

**d. Select query with JOIN**

postgres/postgres@PostgreSQL 13

Query Editor Query History

```
1 explain analyse select payment.amount, buys.customer_email from payment
2 inner join buys on payment.customer_email = buys.customer_email;
```

Data Output Explain Messages Notifications

QUERY PLAN

text

1	Hash Join (cost=18.77..88.53 rows=1443 width=82) (actual time=1.155..1.159 rows=3 loops=1)
2	Hash Cond: ((buys.customer_email)::text = (payment.customer_email)::text)
3	-> Seq Scan on buys (cost=0.00..17.40 rows=740 width=78) (actual time=0.554..0.555 rows=3 loops=1)
4	-> Hash (cost=13.90..13.90 rows=390 width=82) (actual time=0.545..0.546 rows=3 loops=1)
5	Buckets: 1024 Batches: 1 Memory Usage: 9kB
6	-> Seq Scan on payment (cost=0.00..13.90 rows=390 width=82) (actual time=0.157..0.160 rows=3 loops=1)
7	Planning Time: 23.755 ms
8	Execution Time: 1.205 ms

**e. Select query with aggregation**



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

postgres/postgres@PostgreSQL 13

Query Editor Query History

```
1 explain analyse select max(cart_total), min(cart_total), avg(cart_total) from cart;
```

Data Output Explain Messages Notifications

	QUERY PLAN	
	text	
1	Aggregate (cost=22.43..22.44 rows=1 width=40) (actual time=0.033..0.034 rows=1 loops=1)	
2	-> Seq Scan on cart (cost=0.00..17.10 rows=710 width=4) (actual time=0.022..0.023 rows=3 loops=1)	
3	Planning Time: 0.146 ms	
4	Execution Time: 0.072 ms	

Other:

postgres/postgres@PostgreSQL 13

Query Editor Query History

```
1 explain analyse select customer_email, cart_id from payment
2 intersect
3 select customer_email, cart_id from cart;
```

Data Output Explain Messages Notifications

	QUERY PLAN	
	text	
1	HashSetOp Intersect (cost=0.00..53.00 rows=200 width=86) (actual time=0.057..0.060 rows=3 loops=1)	
2	-> Append (cost=0.00..47.50 rows=1100 width=86) (actual time=0.029..0.048 rows=6 loops=1)	
3	-> Subquery Scan on "*SELECT* 1" (cost=0.00..17.80 rows=390 width=86) (actual time=0.029..0.031 rows=3 loops=1)	
4	-> Seq Scan on payment (cost=0.00..13.90 rows=390 width=82) (actual time=0.026..0.028 rows=3 loops=1)	
5	-> Subquery Scan on "*SELECT* 2" (cost=0.00..24.20 rows=710 width=86) (actual time=0.014..0.015 rows=3 loops=1)	
6	-> Seq Scan on cart (cost=0.00..17.10 rows=710 width=82) (actual time=0.014..0.014 rows=3 loops=1)	
7	Planning Time: 0.232 ms	
8	Execution Time: 0.102 ms	



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

postgres/postgres@PostgreSQL 13

Query Editor Query History

```
1 explain analyse select first_name from customer, customer_address
2 where email_id = customer_email and
3     std_code like '____77';
```

Data Output Explain Messages Notifications

QUERY PLAN

1	Hash Join (cost=13.78..26.94 rows=12 width=48) (actual time=0.144..0.147 rows=1 loops=1)
2	Hash Cond: ((customer.email_id)::text = (customer_address.customer_email)::text)
3	-> Seq Scan on customer (cost=0.00..12.50 rows=250 width=126) (actual time=0.034..0.035 rows=6 loops=1)
4	-> Hash (cost=13.63..13.63 rows=12 width=78) (actual time=0.096..0.097 rows=1 loops=1)
5	Buckets: 1024 Batches: 1 Memory Usage: 9kB
6	-> Seq Scan on customer_address (cost=0.00..13.63 rows=12 width=78) (actual time=0.092..0.092 rows=1 loops=1)
7	Filter: (std_code ~~ '____77':text)
8	Rows Removed by Filter: 2
9	Planning Time: 23.584 ms
10	Execution Time: 0.182 ms

### Indexing:

postgres/postgres@PostgreSQL 13

Query Editor Query History

```
1 explain analyse select * from products where discount=10;
```

Data Output Explain Messages Notifications

QUERY PLAN

1	Seq Scan on products (cost=0.00..13.00 rows=1 width=312) (actual time=0.024..0.026 rows=1 loops=1)
2	Filter: (discount = 10)
3	Rows Removed by Filter: 4
4	Planning Time: 0.131 ms
5	Execution Time: 0.051 ms





**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

postgres/postgres@PostgreSQL 13

Query Editor Query History

```
1 EXPLAIN SELECT * FROM products WHERE discount = 10;  
2  
3 CREATE INDEX indexproductsdiscount ON products(discount);  
4  
5 EXPLAIN SELECT * FROM products WHERE discount = 10;  
6  
7 DROP INDEX indexproductsdiscount;  
8
```

Data Output Explain Messages Notifications

CREATE INDEX

Query returned successfully in 101 msec.

postgres/postgres@PostgreSQL 13

Query Editor Query History

```
1 EXPLAIN SELECT * FROM products WHERE discount = 10;  
2  
3 CREATE INDEX indexproductsdiscount ON products(discount);  
4  
5 EXPLAIN SELECT * FROM products WHERE discount = 10;  
6  
7 DROP INDEX indexproductsdiscount;  
8
```

Data Output Explain Messages Notifications

DROP INDEX

Query returned successfully in 103 msec.



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

## Implement

- Indexing on foreign key
- Indexing on composite key
- Indexing on secondary key

### CREATING INDEX

CREATE INDEX SSNO ON SOLDIERS(SSNO DESC);

```
392 //CREATING INDEX
393
394 CREATE INDEX SSNO ON SOLDIERS(SSNO DESC);
395
```

Data Output Explain Messages Notifications

CREATE INDEX

Query returned successfully in 450 msec.

EXPLAIN SELECT \* FROM SOLDIERS WHERE SSNO >604 order by NAME;

```
437
438 EXPLAIN SELECT * FROM SOLDIERS WHERE SSNO >604 order by NAME;
439
440
```

Data Output Explain Messages Notifications

QUERY PLAN  
text

- |   |  |
|---|--|
| 1 | Sort (cost=1.19..1.20 rows=4 width=306)                          |
| 2 | [...] Sort Key: name   |
| 3 | [...] -> Seq Scan on soldiers (cost=0.00..1.15 rows=4 width=306) |
| 4 | [...] Filter: (ssno > 604)                                       |



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

### Simple select query

EXPLAIN (ANALYZE) SELECT \* FROM SOLDIERS;

```
438 EXPLAIN (ANALYZE) SELECT * FROM SOLDIERS;
```

439	
	Data Output Explain Messages Notifications
	<b>QUERY PLAN</b>
	text
1	Seq Scan on soldiers (cost=0.00..1.12 rows=12 width=306) (actual time=0.575..0.577 rows=12 loops=1)
2	Planning Time: 146.001 ms
3	Execution Time: 0.590 ms

### Select query with query

EXPLAIN ANALYZE SELECT \* FROM SOLDIERS WHERE NAME = 'YOGENDRA SINGH';

```
402
403 //EXPLAIN ANALYZER
404
405 EXPLAIN ANALYZE SELECT * FROM SOLDIERS WHERE NAME = 'YOGENDRA SINGH';
```

	Data Output Explain Messages Notifications
	<b>QUERY PLAN</b>
	text
1	Seq Scan on soldiers (cost=0.00..1.15 rows=1 width=306) (actual t...
2	[...] Filter: ((name)::text = 'YOGENDRA SINGH':text)
3	[...] Rows Removed by Filter: 11
4	Planning Time: 0.089 ms
5	Execution Time: 0.023 ms

### Select query with order by query

EXPLAIN ANALYZE SELECT \* FROM SOLDIERS WHERE SSNO >604 order by NAME;

```
407 EXPLAIN ANALYZE SELECT * FROM SOLDIERS WHERE SSNO >604 order by NAME;
```

408	
409	
	Data Output Explain Messages Notifications
	<b>QUERY PLAN</b>
	text
1	Sort (cost=1.19..1.20 rows=4 width=306) (actual time=0.051..0.052 rows=8 loops=1)
2	[...] Sort Key: name
3	[...] Sort Method: quicksort Memory: 26kB
4	[...] Seq Scan on soldiers (cost=0.00..1.15 rows=4 width=306) (actual time=0.009..0.010 rows=8 loops=1)
5	[...] Filter: (ssno > 604)
6	[...] Rows Removed by Filter: 4



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

### Select query with join

EXPLAIN ANALYZE SELECT \*FROM RESOURCES  
RIGHT JOIN ITEMS  
ON resources.rname = items.ideptname;

```
409
410 EXPLAIN ANALYZE SELECT *FROM RESOURCES
411 RIGHT JOIN ITEMS
412 ON resources.rname = items.ideptname;
413
```

Data Output	Explain	Messages	Notifications
<b>QUERY PLAN</b>			
text			
1	Hash Left Join (cost=15.85..29.02 rows=250 width=568) (actual time=9.653..9.663 rows=9 loops=1)		
2	[...] Hash Cond: (items.ideptname = resources.rname)		
3	[...] -> Seq Scan on items (cost=0.00..12.50 rows=250 width=288) (actual time=0.025..0.027 rows=9 loops=1)		
4	[...] -> Hash (cost=12.60..12.60 rows=260 width=280) (actual time=9.595..9.596 rows=5 loops=1)		
5	[...] Buckets: 1024 Batches: 1 Memory Usage: 9kB		
6	[...] -> Seq Scan on resources (cost=0.00..12.60 rows=260 width=280) (actual time=0.009..0.019 rows=5 loops=1)		
7	Planning Time: 30.914 ms		
8	Execution Time: 9.751 ms		

### Select query with aggregation

EXPLAIN ANALYZE SELECT MAX(SALARY) , MIN(SALARY) , AVG(SALARY) FROM SOLDIERS

```
415
416 //Select query with aggregation
417 EXPLAIN ANALYZE SELECT MAX(SALARY) , MIN(SALARY) , AVG(SALARY) FROM SOLDIERS
418
419
420
421
```

Data Output	Explain	Messages	Notifications
<b>QUERY PLAN</b>			
text			
1	Aggregate (cost=1.21..1.22 rows=1 width=40) (actual time=6.785..6.786 rows=1 loops=1)		
2	[...] -> Seq Scan on soldiers (cost=0.00..1.12 rows=12 width=4) (actual time=0.010..0.014 rows=12 loops=1)		
3	Planning Time: 0.162 ms		
4	Execution Time: 6.815 ms		



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

### Indexing on primary key

EXPLAIN ANALYZE SELECT SSNO FROM SOLDIERS

418	
419	//PRIMARY KEY
420	EXPLAIN ANALYZE SELECT SSNO FROM SOLDIERS
421	
Data Output Explain Messages Notifications	
QUERY PLAN	
text	
1	Seq Scan on soldiers (cost=0.00..1.12 rows=12 width=4) (actual time=0.009..0.010 rows=12 loops=1)
2	Planning Time: 0.074 ms
3	Execution Time: 0.018 ms

### Indexing on secondary key

EXPLAIN ANALYZE SELECT NAME FROM SOLDIERS

422	//SECONDARY KEY
423	EXPLAIN ANALYZE SELECT NAME FROM SOLDIERS
424	
Data Output Explain Messages Notifications	
QUERY PLAN	
text	
1	Seq Scan on soldiers (cost=0.00..1.12 rows=12 width=118) (actual time=0.014..0.015 rows=12 loops=1)
2	Planning Time: 0.077 ms
3	Execution Time: 0.028 ms

### Indexing on foreign key

CREATE INDEX INDEX2 ON DEPARTMENT(DNAME)

EXPLAIN ANALYZE SELECT DNAME FROM DEPARTMENT

420	
429	//FOREIGN KEY
430	CREATE INDEX INDEX2 ON DEPARTMENT(DNAME)
431	EXPLAIN ANALYZE SELECT DNAME FROM DEPARTMENT
432	
Data Output Explain Messages Notifications	
QUERY PLAN	
text	
1	Seq Scan on department (cost=0.00..13.30 rows=330 width=104) (actual time=0.028..0.029 rows=10 loops=1)
2	Planning Time: 21.575 ms
3	Execution Time: 0.042 ms



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

### Indexing on composite key

CREATE INDEX INDEX4 ON DEPARTMENT(DNAME,DRNO)

EXPLAIN ANALYZE SELECT DNAME,DRNO FROM DEPARTMENT

433	//COMPOSITE KEY
434	CREATE INDEX INDEX4 ON DEPARTMENT(DNAME,DRNO)
435	EXPLAIN ANALYZE SELECT DNAME,DRNO FROM DEPARTMENT

Data Output	Explain	Messages	Notifications
QUERY PLAN			
text			
1	Seq Scan on department (cost=0.00..1.10 rows=10 width=108) (actual time=0.013..0.015 rows=10 loops=1)		
2	Planning Time: 1.524 ms		
3	Execution Time: 0.027 ms		

### Post Lab Question:

**Q1. Illustrate with an example Heuristic based query optimization with suitable example**

**Ans.**

Heuristic based optimization uses rule-based optimization approaches for query optimization. These algorithms have polynomial time and space complexity, which is lower than the exponential complexity of exhaustive search-based algorithms.

Some of the common rules are –

- Perform select and project operations before join operations. This is done by moving the select and project operations down the query tree. This reduces the number of tuples available for join.
- Perform the most restrictive select/project operations at first before the other operations.
- Avoid cross-product operation since they result in very large-sized intermediate tables.



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

Process for heuristics optimization

1. The parser of a high-level query generates an initial internal representation;
2. Apply heuristics rules to optimize the internal representation.
3. A query execution plan is generated to execute groups of operations based on the access paths available on the files involved in the query.
4. The main heuristic is to apply first the operations that reduce the size of intermediate results.

E.g., Apply SELECT and PROJECT operations before applying the JOIN or other binary operations.

Query tree:

- A tree data structure that corresponds to a relational algebra expression. It represents the input relations of the query as leaf nodes of the tree, and represents the relational algebra operations as internal nodes.

Query graph:

- A graph data structure that corresponds to a relational calculus expression. It does not indicate an order on which operations to perform first. There is only a single graph corresponding to each query.

- SQL query:

SELECT LNAME, FNAME FROM EMPLOYEE WHERE SALARY < (SELECT MAX (SALARY) FROM EMPLOYEE WHERE DNO=5);

(SELECT MAX (SALARY) FROM EMPLOYEE WHERE DNO=5)

MAXSALARY ( $\sigma_{DNO = 5}(EMPLOYEE)$ )

The Outer block SELECT LNAME;

FNAME FROM EMPLOYEE WHERE SALARY < C



**K. J. Somaiya College of Engineering, Mumbai-77**  
(Autonomous College Affiliated to University of Mumbai)

**Q2. List different algorithms of Select, Project , Join with cost**

**Ans.**

**1. Select :-**

- 1) Linear search
- 2) Binary search
- 3) Using a primary index
- 4) Using a primary index to retrieve multiple records: The comparison condition is  $>$ ,  $\leq$ ,  $<$ ,  $\geq$  on a key field with a primary index .
- 5) Using a clustering index to retrieve multiple records: The selection condition involves an equality comparison on a non-key attribute with a clustering index.

**2. Project:-**

- a. CARTESIAN PRODUCT operation
- b. UNION operations
- c. INTERSECTION operations
- d. DIFFERENCE operations

**3. Join :-**

- a. Nested loop join.
- b. Block nested loop join.
- c. Single-loop join
- d. Sort-merge join
- e. Hash-join